

Attribution 4.0 International (CC BY 4.0)

<https://creativecommons.org/licenses/by/4.0/> (<https://creativecommons.org/licenses/by/4.0/>).

[https://en.wikipedia.org/wiki/S.M.A.R.T.#Known\\_ATA\\_S.M.A.R.T.\\_attributes](https://en.wikipedia.org/wiki/S.M.A.R.T.#Known_ATA_S.M.A.R.T._attributes)  
([https://en.wikipedia.org/wiki/S.M.A.R.T.#Known\\_ATA\\_S.M.A.R.T.\\_attributes](https://en.wikipedia.org/wiki/S.M.A.R.T.#Known_ATA_S.M.A.R.T._attributes)).

```
In [1]: import pandas
import sys
print(sys.version_info)
print('pandas',pandas.__version__)
import glob
import pickle
import numpy
import time
import matplotlib.pyplot as plt

sys.version_info(major=3, minor=6, micro=6, releaselevel='final', serial=0)
pandas 0.23.4

In [2]: df_header_only=pandas.read_csv('zipped_data/data_Q2_2018.zip_folder/2018
-04-01.csv',nrows=3)
nonsmart_cols=[]
for colname in df_header_only.columns:
    if 'smart_' not in colname:
        nonsmart_cols.append(colname)

In [3]: nonsmart_cols.append('smart_241_raw') # written
nonsmart_cols.append('smart_242_raw') # read
nonsmart_cols.append('smart_9_raw') # power-on hours
nonsmart_cols.remove('capacity_bytes')

In [4]: list_of_csvs = glob.glob('zipped_data/**/*.csv', recursive=True)
len(list_of_csvs)

Out[4]: 2092

In [5]: start_time=time.time()
list_of_df=[]
for csv_file in list_of_csvs:
    df=pandas.read_csv(csv_file,nrows=2)
    if 'smart_241_raw' in df.columns:
        df=pandas.read_csv(csv_file,usecols=nonsmart_cols)
        df = df[df['failure']==1]
        list_of_df.append(df)
print('elapsed:',time.time()-start_time,'seconds')

elapsed: 446.1708550453186 seconds
```

```
In [6]: df = pandas.concat(list_of_df)
print(df.shape)
#df.dropna(how='any',inplace=True)
#print(df.shape)
#df.head()
```

(8743, 7)

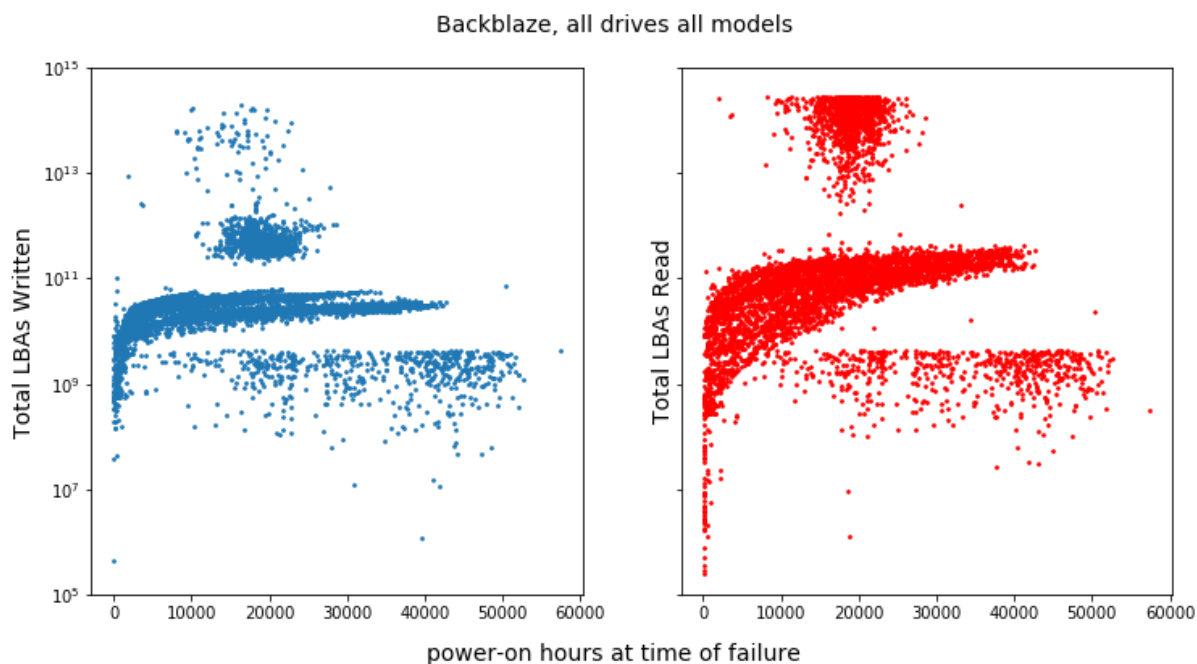
## Logical Block Addresses read/written versus power-on hours

```
In [59]: def lba_plot_vs_poh(y_lower,y_upper,log_bool):
f, (ax1, ax2) = plt.subplots(1, 2, sharey=True,figsize=(12, 6))
ax1.scatter(x=df['smart_9_raw'],y=df['smart_241_raw'],s=3)
ax1.set_ylabel('Total LBAs Written',fontsize=14)
ax1.set_yscale('log')

ax2.scatter(x=df['smart_9_raw'],y=df['smart_242_raw'],color='r',s=3)
ax2.set_ylabel('Total LBAs Read',fontsize=14);
plt.ylim([y_lower,y_upper])
if log_bool:
    ax2.set_yscale('log')

f.text(0.5, 0.04, 'power-on hours at time of failure', ha='center',
va='center',fontsize=14);
f.text(0.5, 0.94, 'Backblaze, all drives all models', ha='center', v
a='center',fontsize=14);
return
```

```
In [60]: lba_plot_vs_poh(y_lower=100000,y_upper=1000000000000000,log_bool=True)
```



## zoom in to the "low LBA read/written" range of values

```
lba_plot_vs_poh(y_lower=0,y_upper=100000000000,log_bool=True)
```

## zoom out to the "medium LBA read/written" range of values

```
lba_plot_vs_poh(y_lower=0,y_upper=100000000000,log_bool=False)
```

## zoom out again to the "high LBA read/written" range of values

```
lba_plot_vs_poh(y_lower=0,y_upper=1000000000000,log_bool=False)
```

## max range for y-axis

```
lba_plot_vs_poh(y_lower=0,y_upper=1E15,log_bool=True)
```

## per model

```
def make_plot_per_model(drive_model): if not (df[df['model']==drive_model]['smart_241_raw'].isnull().all()): f, (ax1,
ax2) = plt.subplots(1, 2, sharey=True,figsize=(12, 6)) ax1.scatter(x=df[df['model']==drive_model]
['smart_9_raw'],y=df[df['model']==drive_model]['smart_241_raw'], s=3,label=drive_model) ax1.set_ylabel('Total
Logical Block Addresses Written',fontsize=14) ax1.set_yscale('log') ax2.scatter(x=df[df['model']==drive_model]
['smart_9_raw'],y=df[df['model']==drive_model]['smart_242_raw'],color='r', s=3,label=drive_model)
ax2.set_ylabel('Total Logical Block Addresses Read',fontsize=14); plt.ylim([100000,1000000000000000])
ax2.set_yscale('log') f.text(0.5, 0.04, 'power-on hours at time of removal', ha='center', va='center',fontsize=14);
f.text(0.5, 0.94, 'Backblaze: '+drive_model, ha='center', va='center',fontsize=14);drive_model='ST4000DM000'
print(df[df['model']==drive_model].shape) make_plot_per_model(drive_model)ser = df['model'].value_counts() for
drive_model in ser[ser>100].index: # only show results if there are more than 100 instances of that drive model
being removed make_plot_per_model(drive_model)
```

```

In [61]: # LBA*(4*1012) bytes to tb

def make_TB_plot_per_model(drive_model):

    if not (df[df['model']==drive_model]['smart_241_raw'].isnull().all
    ()):
        f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(12, 6))
        ax1.scatter(x=df[df['model']==drive_model]['smart_9_raw']/24/30,
y=df[df['model']==drive_model]['smart_241_raw']*(4*1012)/(1E12),
                    s=3, label=drive_model)
        ax1.set_ylabel('Total terabytes Written', fontsize=14)
        ax1.set_yscale('log')


        ax2.scatter(x=df[df['model']==drive_model]['smart_9_raw']/24/30,
y=df[df['model']==drive_model]['smart_242_raw']*(4*1012)/(1E12), color=
'r',
                    s=3, label=drive_model)
        ax2.set_ylabel('Total terabytes Read', fontsize=14);
        plt.ylim([0.1, 100000])
        ax2.set_yscale('log')

        f.text(0.5, 0.04, 'power-on months at time of removal', ha='cent
er', va='center', fontsize=14);
        f.text(0.5, 0.94, 'Backblaze: '+drive_model, ha='center', va='ce
nter', fontsize=14);

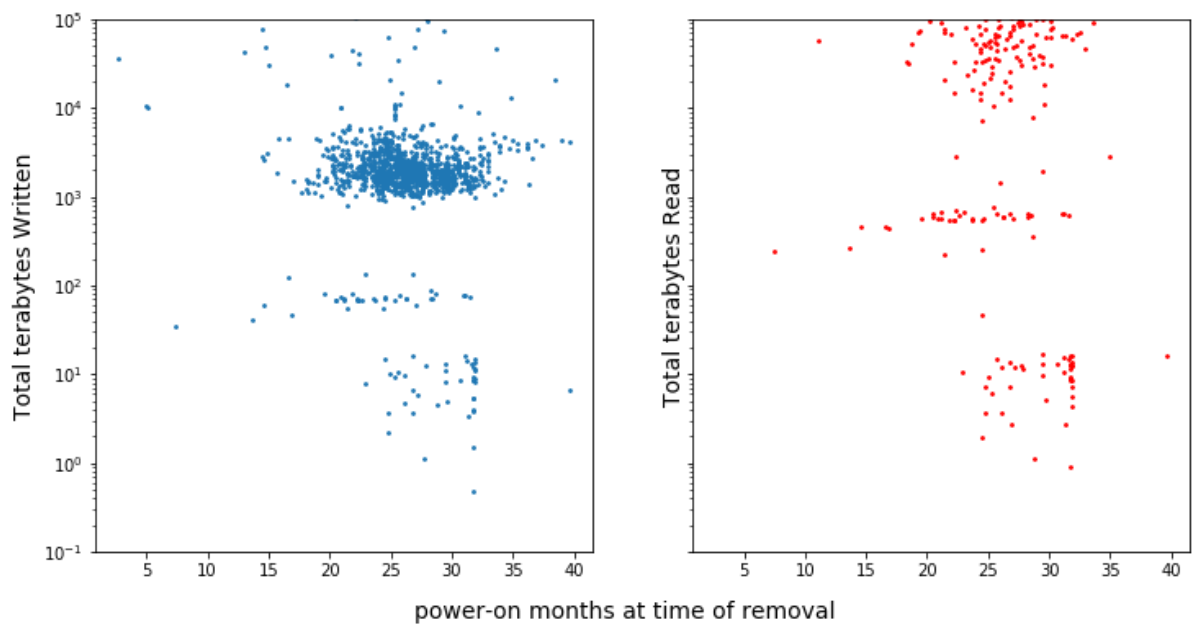
```

```
In [62]: ser = df['model'].value_counts()
for drive_model in ser[ser>100].index: # only show results if there are
    more than 100 instances of that drive model being removed
    make_TB_plot_per_model(drive_model)
```

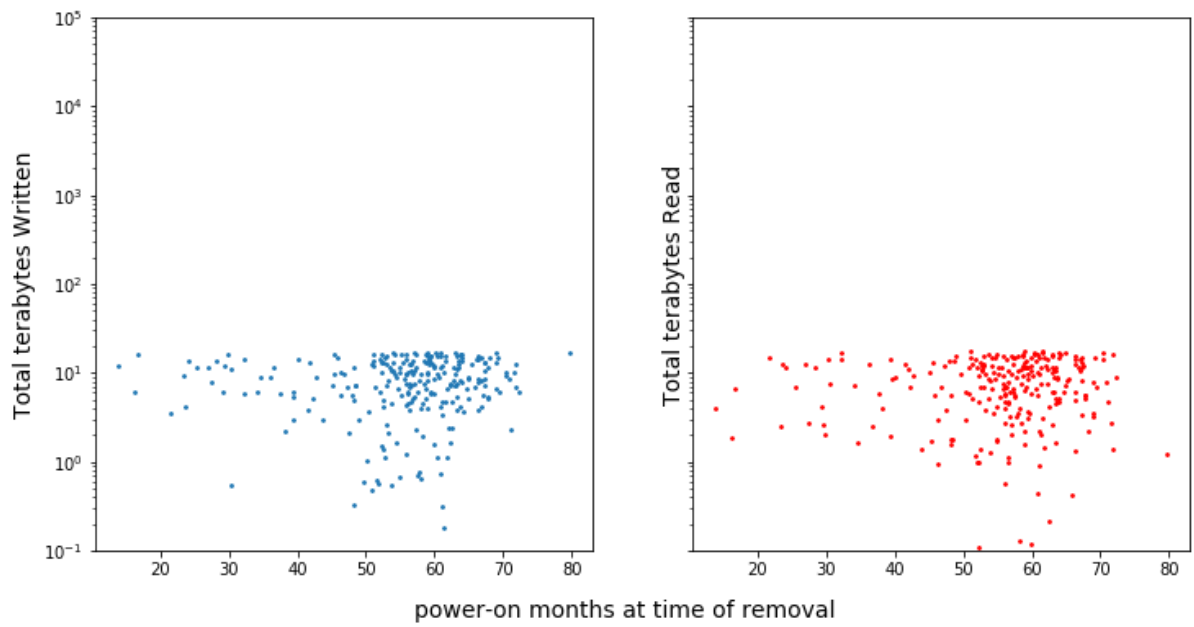
Backblaze: ST4000DM000




Backblaze: ST3000DM001




Backblaze: ST31500541AS



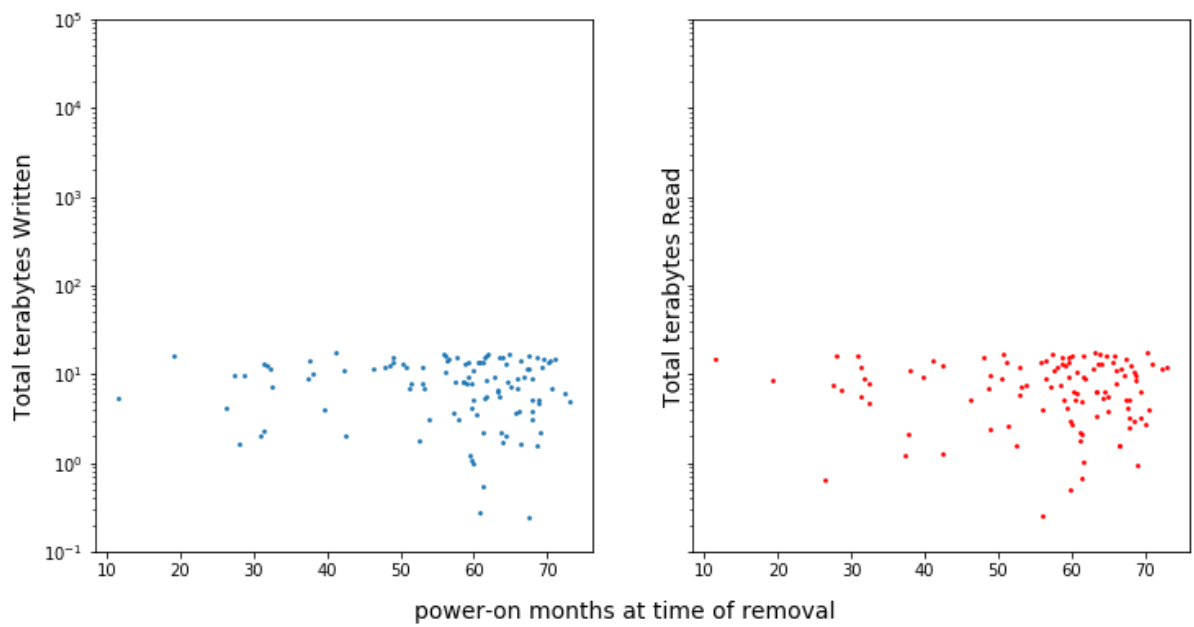
Backblaze: ST12000NM0007



Backblaze: ST8000DM002



Backblaze: ST31500341AS





## Backblaze: ST8000NM0055

